

4 Classification

By: Udit (based on ISLR)

Setup

LDA & QDA is part of **MASS** library.

Naive Bayes is part of **e1071** library.

KNN is part of **class** library.

```
library(ISLR2)
library(MASS)

## 
## Attaching package: 'MASS'

## The following object is masked from 'package:ISLR2':
## 
##     Boston

library(class)
library(e1071)

# Stock Market data
attach(Smarket)
dim(Smarket)

## [1] 1250    9

names(Smarket)

## [1] "Year"      "Lag1"       "Lag2"       "Lag3"       "Lag4"       "Lag5"
## [7] "Volume"    "Today"      "Direction"

summary(Smarket)

##      Year          Lag1          Lag2          Lag3
##  Min. :2001  Min. :-4.922000  Min. :-4.922000  Min. :-4.922000
##  1st Qu.:2002 1st Qu.:-0.639500 1st Qu.:-0.639500 1st Qu.:-0.640000
##  Median :2003 Median : 0.039000 Median : 0.039000 Median : 0.038500
##  Mean   :2003 Mean  : 0.003834 Mean  : 0.003919 Mean  : 0.001716
##  3rd Qu.:2004 3rd Qu.: 0.596750 3rd Qu.: 0.596750 3rd Qu.: 0.596750
##  Max.   :2005 Max.  : 5.733000 Max.  : 5.733000 Max.  : 5.733000
##      Lag4          Lag5          Volume        Today
##  Min. :-4.922000  Min. :-4.922000  Min. :0.3561  Min. :-4.922000
```

```

## 1st Qu.:-0.640000 1st Qu.:-0.64000 1st Qu.:1.2574 1st Qu.:-0.639500
## Median : 0.038500 Median : 0.03850 Median :1.4229 Median : 0.038500
## Mean : 0.001636 Mean : 0.00561 Mean :1.4783 Mean : 0.003138
## 3rd Qu.: 0.596750 3rd Qu.: 0.59700 3rd Qu.:1.6417 3rd Qu.: 0.596750
## Max. : 5.733000 Max. : 5.73300 Max. :3.1525 Max. : 5.733000
## Direction
## Down:602
## Up :648
##
##
##
##
```

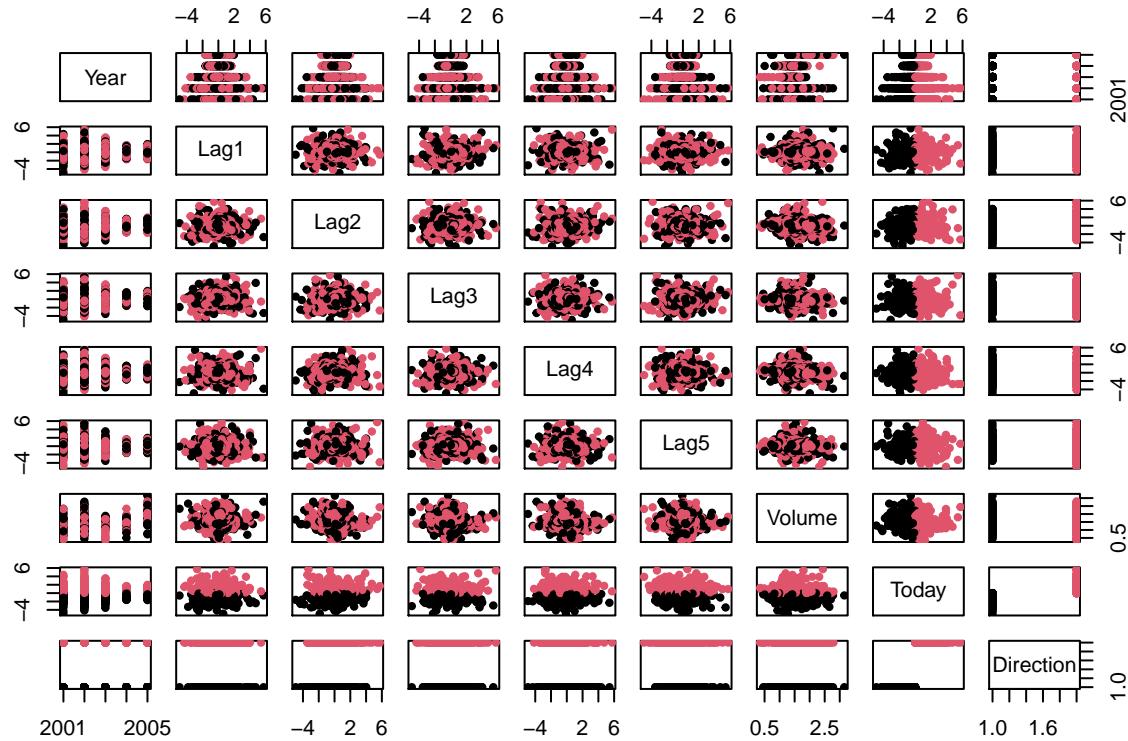
```
round(cor(Smarket[,-9]),2)
```

```

##      Year Lag1 Lag2 Lag3 Lag4 Lag5 Volume Today
## Year  1.00  0.03  0.03  0.03  0.04  0.03   0.54  0.03
## Lag1  0.03  1.00 -0.03 -0.01  0.00 -0.01   0.04 -0.03
## Lag2  0.03 -0.03  1.00 -0.03 -0.01  0.00   -0.04 -0.01
## Lag3  0.03 -0.01 -0.03  1.00 -0.02 -0.02   -0.04  0.00
## Lag4  0.04  0.00 -0.01 -0.02  1.00 -0.03   -0.05 -0.01
## Lag5  0.03 -0.01  0.00 -0.02 -0.03  1.00   -0.02 -0.03
## Volume 0.54  0.04 -0.04 -0.04 -0.05 -0.02   1.00  0.01
## Today  0.03 -0.03 -0.01  0.00 -0.01 -0.03   0.01  1.00

```

```
pairs(Smarket, col=Smarket$Direction, pch=20)
```



Logistic Regression

```
glm.fit = glm(Direction~.-Year-Today, data=Smarket, family=binomial)
summary(glm.fit)
```

```
##  
## Call:  
## glm(formula = Direction ~ . - Year - Today, family = binomial,  
##       data = Smarket)  
##  
## Deviance Residuals:  
##    Min      1Q  Median      3Q     Max  
## -1.446  -1.203   1.065   1.145   1.326  
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept) -0.126000  0.240736 -0.523   0.601  
## Lag1        -0.073074  0.050167 -1.457   0.145  
## Lag2        -0.042301  0.050086 -0.845   0.398  
## Lag3         0.011085  0.049939  0.222   0.824  
## Lag4         0.009359  0.049974  0.187   0.851  
## Lag5         0.010313  0.049511  0.208   0.835  
## Volume       0.135441  0.158360  0.855   0.392  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
## Null deviance: 1731.2 on 1249 degrees of freedom  
## Residual deviance: 1727.6 on 1243 degrees of freedom  
## AIC: 1741.6  
##  
## Number of Fisher Scoring iterations: 3
```

```
# Predicted probabilities are close to 50% as expected
predict(glm.fit, type="response")[1:5]
```

```
##          1          2          3          4          5
## 0.5070841 0.4814679 0.4811388 0.5152224 0.5107812
```

```
# Probability to Classification
glm.probs = predict(glm.fit, type="response")
glm.pred = ifelse(glm.probs>0.5, "Up", "Down")
```



```
# Confusion Matrix
attach(Smarket)
```

```
## The following objects are masked from Smarket (pos = 3):
##  
##      Direction, Lag1, Lag2, Lag3, Lag4, Lag5, Today, Volume, Year
```

```

table(glm.pred, Direction)

##          Direction
## glm.pred Down Up
##      Down 145 141
##      Up   457 507

mean(glm.pred==Direction)

## [1] 0.5216

# Make Training & Test data
train = Year<2005
glm.fit = glm(Direction~. - Today -Year, data=Smarket, family=binomial,
subset=train)
glm.probs = predict(glm.fit, newdata=Smarket[!train,], type="response")
glm preds = ifelse(glm.probs>0.5, "Up", "Down")

# Check performance - possible overfitting
Direction.test = Smarket$Direction[!train]
table(glm preds, Direction.test)

##          Direction.test
## glm preds Down Up
##      Down 77 97
##      Up   34 44

mean(glm preds==Direction.test)

## [1] 0.4801587

# Fit smaller model
glm.fit = glm(Direction~Lag1+Lag2, data=Smarket, family=binomial, subset=train)
glm.probs = predict(glm.fit, newdata=Smarket[!train,], type="response")
contrasts(Direction)

##      Up
## Down 0
## Up   1

glm preds = ifelse(glm.probs>0.5, "Up", "Down")
table(glm preds, Direction.test)

##          Direction.test
## glm preds Down Up
##      Down 35 35
##      Up   76 106

```

```

mean(glm preds==Direction.test)

## [1] 0.5595238

106/(106+76)

## [1] 0.5824176

summary(glm.fit)

##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2, family = binomial, data = Smarket,
##      subset = train)
##
## Deviance Residuals:
##    Min      1Q  Median      3Q     Max
## -1.345  -1.188   1.074   1.164   1.326
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.03222   0.06338   0.508   0.611
## Lag1        -0.05562   0.05171  -1.076   0.282
## Lag2        -0.04449   0.05166  -0.861   0.389
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1383.3 on 997 degrees of freedom
## Residual deviance: 1381.4 on 995 degrees of freedom
## AIC: 1387.4
##
## Number of Fisher Scoring iterations: 3

```

Linear Discriminant Analysis

```

lda.fit = lda(Direction~Lag1+Lag2, subset=Year<2005)
lda.fit

##
## Call:
## lda(Direction ~ Lag1 + Lag2, subset = Year < 2005)
##
## Prior probabilities of groups:
##       Down      Up
## 0.491984 0.508016
##
## Group means:
##           Lag1      Lag2
## Down  0.04279022 0.03389409
## Up   -0.03954635 -0.03132544

```

```

##  

## Coefficients of linear discriminants:  

##           LD1  

## Lag1 -0.6420190  

## Lag2 -0.5135293  

Smarket.test = subset(Smarket, Year>=2005)  

lda.pred = predict(lda.fit, Smarket.test)  

data.frame(lda.pred)[1:5,]  

##      class posterior.Down posterior.Up       LD1  

## 999     Up      0.4901792   0.5098208  0.08293096  

## 1000    Up      0.4792185   0.5207815  0.59114102  

## 1001    Up      0.4668185   0.5331815  1.16723063  

## 1002    Up      0.4740011   0.5259989  0.83335022  

## 1003    Up      0.4927877   0.5072123 -0.03792892  

table(lda.pred$class, Smarket.test$Direction)  

##  

##      Down  Up  

##  Down  35  35  

##  Up    76 106  

mean(lda.pred$class == Smarket.test$Direction)  

## [1] 0.5595238

```

Quadratic Discriminant Analysis

```

qda.fit = qda(Direction ~ Lag1 + Lag2, data=Smarket, subset=train)  

qda.fit  

## Call:  

## qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)  

##  

## Prior probabilities of groups:  

##      Down      Up  

## 0.491984 0.508016  

##  

## Group means:  

##           Lag1      Lag2  

## Down  0.04279022  0.03389409  

## Up    -0.03954635 -0.03132544  

qda.pred = predict(qda.fit, Smarket.test)  

table(qda.pred$class, Smarket.test$Direction)

```

```

##          Down   Up
##  Down    30   20
##  Up     81 121

mean(qda.pred$class == Smarket.test$Direction)

```

```
## [1] 0.5992063
```

Naive Bayes

```

nb.fit = naiveBayes(Direction~Lag1+Lag2, data=Smarket, subset=train)
nb.fit

```

```

##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
##
## A-priori probabilities:
## Y
##      Down       Up
## 0.491984 0.508016
##
## Conditional probabilities:
##      Lag1
## Y      [,1]      [,2]
## Down  0.04279022 1.227446
## Up   -0.03954635 1.231668
##
##      Lag2
## Y      [,1]      [,2]
## Down  0.03389409 1.239191
## Up   -0.03132544 1.220765

```

```
mean(Lag1[train][Direction[train]=="Down"])

```

```
## [1] 0.04279022
```

```
sd(Lag1[train][Direction[train]=="Down"])

```

```
## [1] 1.227446
```

```

nb.pred = predict(nb.fit, Smarket.test)
table(nb.pred, Smarket.test$Direction)

```

```

##
## nb.pred  Down   Up
##  Down    28   20
##  Up     83 121

```

```

mean(nb.pred == Smarket.test$Direction)

## [1] 0.5912698

predict(nb.fit, Smarket.test, type="raw")[1:5,]

##          Down      Up
## [1,] 0.4873164 0.5126836
## [2,] 0.4762492 0.5237508
## [3,] 0.4653377 0.5346623
## [4,] 0.4748652 0.5251348
## [5,] 0.4901890 0.5098110

```

KNN

```

attach(Smarket)

## The following objects are masked from Smarket (pos = 3):
## 
##     Direction, Lag1, Lag2, Lag3, Lag4, Lag5, Today, Volume, Year

## The following objects are masked from Smarket (pos = 4):
## 
##     Direction, Lag1, Lag2, Lag3, Lag4, Lag5, Today, Volume, Year

ls()

##  [1] "Direction.test"    "glm.fit"           "glm.pred"         "glm.preds"
##  [5] "glm.probs"         "lda.fit"           "lda.pred"         "nb.fit"
##  [9] "nb.pred"            "qda.fit"           "qda.pred"         "Smarket.test"
## [13] "train"

xlag = cbind(Lag1, Lag2)
xlag[1:5,]

##          Lag1      Lag2
## [1,] 0.381 -0.192
## [2,] 0.959  0.381
## [3,] 1.032  0.959
## [4,] -0.623  1.032
## [5,] 0.614 -0.623

knn.pred = knn(xlag[train,], xlag[!train,], Direction[train], k=1)
table(knn.pred, Direction[!train])

## 
## knn.pred  Down Up
##     Down   43 58
##     Up    68 83

```

```

mean(knn.pred == Direction[!train])

## [1] 0.5

for(i in 1:10){
knn.pred = knn(xlag[train,], xlag[!train,], Direction[train], k=i)
table(knn.pred,Direction[!train])
print(c(i,mean(knn.pred == Direction[!train])))
}

## [1] 1.0 0.5
## [1] 2.000000 0.515873
## [1] 3.0000000 0.5277778
## [1] 4.000000 0.531746
## [1] 5.000000 0.484127
## [1] 6.0 0.5
## [1] 7.0000000 0.4642857
## [1] 8.0000000 0.4801587
## [1] 9.0000000 0.4960317
## [1] 10.0000000 0.5119048

```

KNN - Example 2

```

attach(Caravan)
dim(Caravan)

## [1] 5822 86

names(Caravan) #85 demographic indicators

## [1] "MOSSTYPE" "MAANTHUI" "MGEMOMV" "MGEMLEEF" "MOSHOOFD" "MGODRK"
## [7] "MGODPR" "MGODOV" "MGODGE" "MRELGE" "MRELSA" "MRELOV"
## [13] "MFALLEEN" "MFGEKIND" "MFWEKIND" "MOPLHOOG" "MOPLMIDD" "MOPLLAAG"
## [19] "MBERHOOG" "MBERZELF" "MBERBOER" "MBERMIDD" "MBERARBG" "MBERARBO"
## [25] "MSKA" "MSKB1" "MSKB2" "MSKC" "MSKD" "MHUUR"
## [31] "MHKOOP" "MAUT1" "MAUT2" "MAUTO" "MZFONDS" "MZPART"
## [37] "MINKM30" "MINK3045" "MINK4575" "MINK7512" "MINK123M" "MINKGEM"
## [43] "MKOOPKLA" "PWAPART" "PWABEDR" "PWALAND" "PPERSAUT" "PBESAUT"
## [49] "PMOTSCO" "PVRAAUT" "PAANHANG" "PTRACTOR" "PWERKT" "PBROM"
## [55] "PLEVEN" "PPERSONG" "PGEZONG" "PWAOREG" "PBRAND" "PZEILPL"
## [61] "PPLEZIER" "PFIETS" "PINBOED" "PYSTAND" "AWAPART" "AWABEDR"
## [67] "AWALAND" "APERSAUT" "ABESAUT" "AMOTSCO" "AVRAAUT" "AAANHANG"
## [73] "ATRACTOR" "AWERKT" "ABROM" "ALEVEN" "APERSONG" "AGEZONG"
## [79] "AWAOREG" "ABRAND" "AZEILPL" "APLEZIER" "AFIETS" "AINBOED"
## [85] "ABYSTAND" "Purchase"

```

```
summary(Purchase)
```

```

##   No Yes
## 5474 348

```

```

round(348/5822,4)*100

## [1] 5.98

standard.X = scale(Caravan[,-86])
var(Caravan[,1:2]); var(standard.X[,1:2])

##          MOSTYPE    MAANTHUI
## MOSTTYPE  165.0378474 -0.2018823
## MAANTHUI -0.2018823  0.1647078

##          MOSTYPE    MAANTHUI
## MOSTTYPE   1.00000000 -0.03872126
## MAANTHUI -0.03872126  1.00000000

test <- 1:1000
test.X = standard.X[test,]
test.Y = Purchase[test]

train.X = standard.X[-test,]
train.Y = Purchase[-test]

set.seed(1)
knn.pred <- knn(train.X, test.X, train.Y, k=1)
mean(test.Y == knn.pred)

## [1] 0.882

# Error rate
mean(test.Y != knn.pred)

## [1] 0.118

mean(test.Y != "No") # no skill, always predicting "no"

## [1] 0.059

# How about only positive cases
table(knn.pred, test.Y)

##      test.Y
## knn.pred No Yes
##       No 873  50
##       Yes  68   9

10/(69+10)

## [1] 0.1265823

```

```
# Changing K (can be tested through CV)
knn.pred <- knn(train.X, test.X, train.Y, k=3)
table(knn.pred, test.Y)
```

```
##           test.Y
## knn.pred  No Yes
##       No  920  54
##      Yes   21   5
```

5/(19+5)

```
## [1] 0.2083333
```

```
knn.pred <- knn(train.X, test.X, train.Y, k=5)
table(knn.pred, test.Y)
```

```
##           test.Y
## knn.pred  No Yes
##       No  930  55
##      Yes   11   4
```

4/15

```
## [1] 0.2666667
```

```
# Compare with Logistic Regression
glm.fits = glm(Purchase ~., data=Caravan, subset=-test, family=binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
glm.probs = predict(glm.fits, Caravan[test,], type="response")
pred.Y = ifelse(glm.probs>0.5, "Yes", "No")
table(pred.Y, test.Y) # Recall = 0
```

```
##           test.Y
## pred.Y  No Yes
##     No  934  59
##    Yes   7   0
```

```
# Changing threshold
pred.Y = ifelse(glm.probs>0.25, "Yes", "No")
table(pred.Y, test.Y) # Recall = 0
```

```
##           test.Y
## pred.Y  No Yes
##     No  919  48
##    Yes   22  11
```

Poisson Regression

```
attach(Bikeshare)
dim(Bikeshare)

## [1] 8645 15

names(Bikeshare)

## [1] "season"      "mnth"        "day"         "hr"          "holiday"
## [6] "weekday"     "workingday"   "weathersit"   "temp"        "atemp"
## [11] "hum"         "windspeed"    "casual"      "registered"  "bikers"

# Linear Regression #####
mod.lm <- lm(bikers~mnth+hr+workingday+temp+weathersit, data=Bikeshare)
summary(mod.lm)

## 
## Call:
## lm(formula = bikers ~ mnth + hr + workingday + temp + weathersit,
##      data = Bikeshare)
## 
## Residuals:
##       Min     1Q Median     3Q    Max 
## -299.00 -45.70  -6.23  41.08 425.29 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -68.632    5.307 -12.932 < 2e-16 ***
## mnthFeb      6.845    4.287   1.597 0.110398  
## mnthMarch    16.551    4.301   3.848 0.000120 ***
## mnthApril    41.425    4.972   8.331 < 2e-16 ***
## mnthMay      72.557    5.641  12.862 < 2e-16 ***
## mnthJune     67.819    6.544  10.364 < 2e-16 ***
## mnthJuly     45.324    7.081   6.401 1.63e-10 ***
## mnthAug      53.243    6.640   8.019 1.21e-15 ***
## mnthSept     66.678    5.925  11.254 < 2e-16 ***
## mnthOct      75.834    4.950  15.319 < 2e-16 ***
## mnthNov      60.310    4.610  13.083 < 2e-16 ***
## mnthDec      46.458    4.271  10.878 < 2e-16 ***
## hr1          -14.579   5.699  -2.558 0.010536 *  
## hr2          -21.579   5.733  -3.764 0.000168 *** 
## hr3          -31.141   5.778  -5.389 7.26e-08 *** 
## hr4          -36.908   5.802  -6.361 2.11e-10 *** 
## hr5          -24.135   5.737  -4.207 2.61e-05 *** 
## hr6          20.600    5.704  3.612 0.000306 *** 
## hr7          120.093   5.693  21.095 < 2e-16 *** 
## hr8          223.662   5.690  39.310 < 2e-16 *** 
## hr9          120.582   5.693  21.182 < 2e-16 *** 
## hr10         83.801    5.705  14.689 < 2e-16 *** 
## hr11         105.423   5.722  18.424 < 2e-16 ***
```

```

## hr12          137.284    5.740  23.916 < 2e-16 ***
## hr13          136.036    5.760  23.617 < 2e-16 ***
## hr14          126.636    5.776  21.923 < 2e-16 ***
## hr15          132.087    5.780  22.852 < 2e-16 ***
## hr16          178.521    5.772  30.927 < 2e-16 ***
## hr17          296.267    5.749  51.537 < 2e-16 ***
## hr18          269.441    5.736  46.976 < 2e-16 ***
## hr19          186.256    5.714  32.596 < 2e-16 ***
## hr20          125.549    5.704  22.012 < 2e-16 ***
## hr21          87.554     5.693  15.378 < 2e-16 ***
## hr22          59.123     5.689  10.392 < 2e-16 ***
## hr23          26.838     5.688   4.719 2.41e-06 ***
## workingday      1.270     1.784   0.711 0.476810
## temp           157.209    10.261  15.321 < 2e-16 ***
## weathersitcloudy/misty -12.890    1.964  -6.562 5.60e-11 ***
## weathersitlight rain/snow -66.494    2.965  -22.425 < 2e-16 ***
## weathersitheavy rain/snow -109.745   76.667  -1.431 0.152341
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ',' 1
##
## Residual standard error: 76.5 on 8605 degrees of freedom
## Multiple R-squared:  0.6745, Adjusted R-squared:  0.6731
## F-statistic: 457.3 on 39 and 8605 DF,  p-value: < 2.2e-16

# Alternate Coding for qualitative variables - no base; Coeff(last level) = Sum of other Coeffs.
contrasts(Bikeshare$hr) = contr.sum(24)
contrasts(Bikeshare$mnth) = contr.sum(12)

mod.lm2 <- lm(bikers~mnth+hr+workingday+temp+weathersit, data=Bikeshare)
summary(mod.lm2)

##
## Call:
## lm(formula = bikers ~ mnth + hr + workingday + temp + weathersit,
##      data = Bikeshare)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -299.00  -45.70   -6.23   41.08  425.29 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 73.5974    5.1322 14.340 < 2e-16 ***
## mnth1      -46.0871    4.0855 -11.281 < 2e-16 ***
## mnth2      -39.2419    3.5391 -11.088 < 2e-16 ***
## mnth3      -29.5357    3.1552  -9.361 < 2e-16 ***
## mnth4      -4.6622    2.7406  -1.701 0.08895 .  
## mnth5       26.4700    2.8508   9.285 < 2e-16 ***
## mnth6       21.7317    3.4651   6.272 3.75e-10 ***
## mnth7      -0.7626    3.9084  -0.195 0.84530  
## mnth8       7.1560    3.5347   2.024 0.04295 *  
## mnth9      20.5912    3.0456   6.761 1.46e-11 *** 
## mnth10     29.7472    2.6995  11.019 < 2e-16 ***
## mnth11     14.2229    2.8604   4.972 6.74e-07 ***

```

```

## hr1          -96.1420   3.9554 -24.307 < 2e-16 ***
## hr2          -110.7213  3.9662 -27.916 < 2e-16 ***
## hr3          -117.7212  4.0165 -29.310 < 2e-16 ***
## hr4          -127.2828  4.0808 -31.191 < 2e-16 ***
## hr5          -133.0495  4.1168 -32.319 < 2e-16 ***
## hr6          -120.2775  4.0370 -29.794 < 2e-16 ***
## hr7          -75.5424   3.9916 -18.925 < 2e-16 ***
## hr8          23.9511   3.9686  6.035  1.65e-09 ***
## hr9          127.5199   3.9500 32.284 < 2e-16 ***
## hr10         24.4399   3.9360  6.209  5.57e-10 ***
## hr11         -12.3407  3.9361 -3.135  0.00172 **
## hr12         9.2814    3.9447  2.353  0.01865 *
## hr13         41.1417   3.9571 10.397 < 2e-16 ***
## hr14         39.8939   3.9750 10.036 < 2e-16 ***
## hr15         30.4940   3.9910  7.641  2.39e-14 ***
## hr16         35.9445   3.9949  8.998 < 2e-16 ***
## hr17         82.3786   3.9883 20.655 < 2e-16 ***
## hr18         200.1249  3.9638 50.488 < 2e-16 ***
## hr19         173.2989  3.9561 43.806 < 2e-16 ***
## hr20         90.1138   3.9400 22.872 < 2e-16 ***
## hr21         29.4071   3.9362  7.471  8.74e-14 ***
## hr22         -8.5883   3.9332 -2.184  0.02902 *
## hr23         -37.0194  3.9344 -9.409 < 2e-16 ***
## workingday    1.2696   1.7845  0.711  0.47681
## temp          157.2094  10.2612 15.321 < 2e-16 ***
## weathersitcloudy/misty -12.8903  1.9643 -6.562  5.60e-11 ***
## weathersitlight rain/snow -66.4944  2.9652 -22.425 < 2e-16 ***
## weathersitheavy rain/snow -109.7446  76.6674 -1.431  0.15234
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 76.5 on 8605 degrees of freedom
## Multiple R-squared:  0.6745, Adjusted R-squared:  0.6731
## F-statistic: 457.3 on 39 and 8605 DF,  p-value: < 2.2e-16

```

```

# No impact to predictions
sum((predict(mod.lm)-predict(mod.lm2))^2)

```

```

## [1] 1.426274e-18

```

```

all.equal(predict(mod.lm), predict(mod.lm2))

```

```

## [1] TRUE

```

```

# Coeff for December
-sum(coef(mod.lm2)[2:12])

```

```

## [1] 0.3705817

```

```

# Plotting Coeff
par(mfrow=c(1,2))
coef.months = c(coef(mod.lm2)[2:12], -sum(coef(mod.lm2)[2:12]))

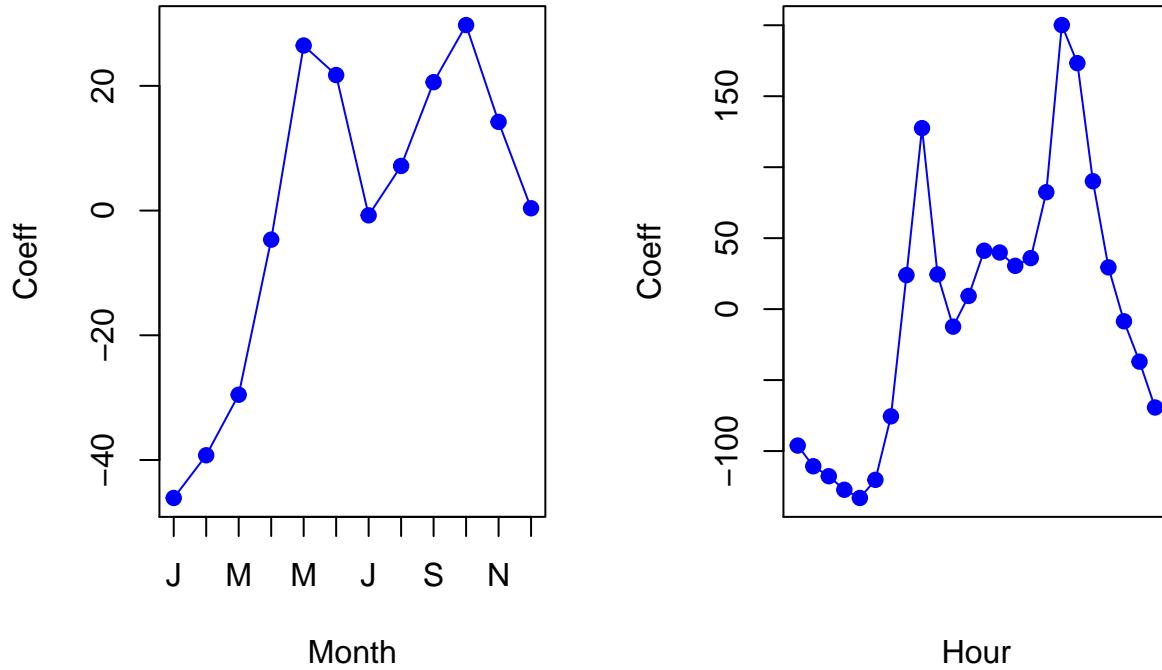
```

```

plot(coef.months, xlab="Month", ylab="Coeff", xaxt="n", col="blue", pch=19, type="o")
axis(side=1, at=1:12, labels=c("J","F","M","A","M","J","J","A","S","O","N","D"))

coef.hrs = c(coef(mod.lm2)[13:35], -sum(coef(mod.lm2)[13:35]))
plot(coef.hrs, xlab="Hour", ylab="Coeff", xaxt="n", col="blue", pch=19, type="o")

```



```

# Poisson Regression #####
mod.pois <- glm(bikers~mnth+hr+workingday+temp+weathersit, data=Bikeshare,
                 family=poisson)
summary(mod.pois)

```

```

##
## Call:
## glm(formula = bikers ~ mnth + hr + workingday + temp + weathersit,
##      family = poisson, data = Bikeshare)
##
## Deviance Residuals:
##      Min        1Q        Median         3Q        Max 
## -20.7574   -3.3441   -0.6549    2.6999   21.9628 
##
## Coefficients:
##                               Estimate Std. Error z value Pr(>|z|)    
## (Intercept)             4.118245   0.006021 683.964 < 2e-16 ***
## mnth1                  -0.670170   0.005907 -113.445 < 2e-16 ***
## mnth2                  -0.444124   0.004860  -91.379 < 2e-16 ***

```

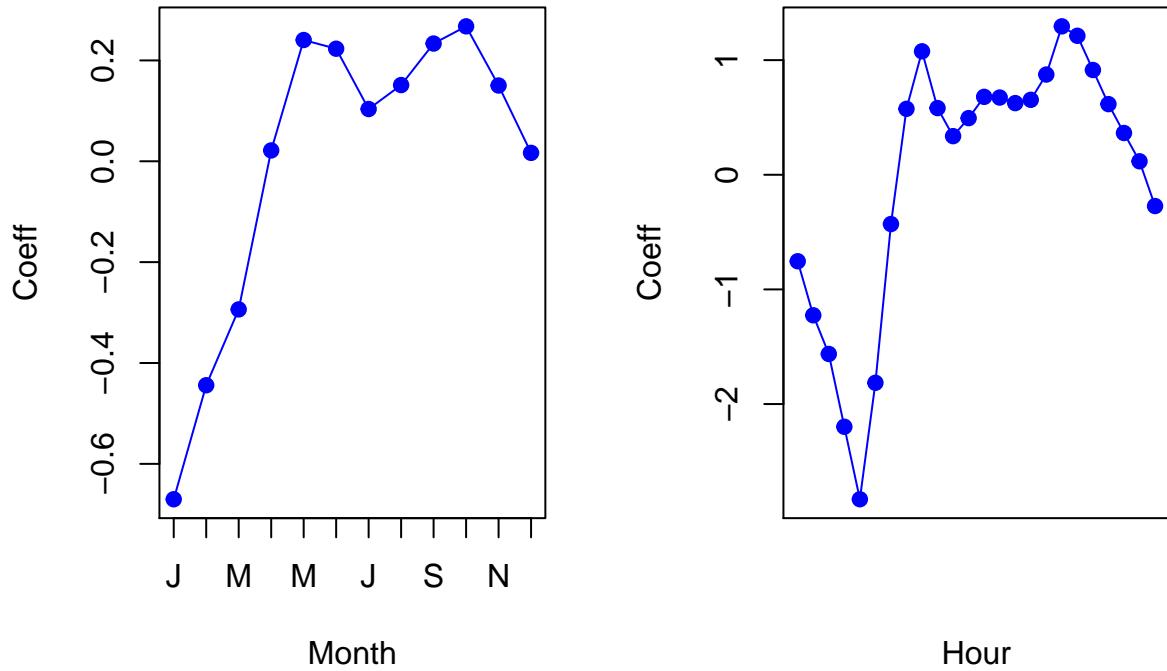
```

## mnth3      -0.293733  0.004144  -70.886 < 2e-16 ***
## mnth4      0.021523  0.003125   6.888 5.66e-12 ***
## mnth5      0.240471  0.002916  82.462 < 2e-16 ***
## mnth6      0.223235  0.003554  62.818 < 2e-16 ***
## mnth7      0.103617  0.004125  25.121 < 2e-16 ***
## mnth8      0.151171  0.003662  41.281 < 2e-16 ***
## mnth9      0.233493  0.003102  75.281 < 2e-16 ***
## mnth10     0.267573  0.002785  96.091 < 2e-16 ***
## mnth11     0.150264  0.003180  47.248 < 2e-16 ***
## hr1        -0.754386  0.007879 -95.744 < 2e-16 ***
## hr2        -1.225979  0.009953 -123.173 < 2e-16 ***
## hr3        -1.563147  0.011869 -131.702 < 2e-16 ***
## hr4        -2.198304  0.016424 -133.846 < 2e-16 ***
## hr5        -2.830484  0.022538 -125.586 < 2e-16 ***
## hr6        -1.814657  0.013464 -134.775 < 2e-16 ***
## hr7        -0.429888  0.006896 -62.341 < 2e-16 ***
## hr8        0.575181  0.004406 130.544 < 2e-16 ***
## hr9        1.076927  0.003563 302.220 < 2e-16 ***
## hr10       0.581769  0.004286 135.727 < 2e-16 ***
## hr11       0.336852  0.004720  71.372 < 2e-16 ***
## hr12       0.494121  0.004392 112.494 < 2e-16 ***
## hr13       0.679642  0.004069 167.040 < 2e-16 ***
## hr14       0.673565  0.004089 164.722 < 2e-16 ***
## hr15       0.624910  0.004178 149.570 < 2e-16 ***
## hr16       0.653763  0.004132 158.205 < 2e-16 ***
## hr17       0.874301  0.003784 231.040 < 2e-16 ***
## hr18       1.294635  0.003254 397.848 < 2e-16 ***
## hr19       1.212281  0.003321 365.084 < 2e-16 ***
## hr20       0.914022  0.003700 247.065 < 2e-16 ***
## hr21       0.616201  0.004191 147.045 < 2e-16 ***
## hr22       0.364181  0.004659  78.173 < 2e-16 ***
## hr23       0.117493  0.005225  22.488 < 2e-16 ***
## workingday  0.014665  0.001955   7.502 6.27e-14 ***
## temp        0.785292  0.011475  68.434 < 2e-16 ***
## weathersitcloudy/misty -0.075231  0.002179 -34.528 < 2e-16 ***
## weathersitlight rain/snow -0.575800  0.004058 -141.905 < 2e-16 ***
## weathersitheavy rain/snow -0.926287  0.166782  -5.554 2.79e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
## Null deviance: 1052921  on 8644  degrees of freedom
## Residual deviance: 228041  on 8605  degrees of freedom
## AIC: 281159
##
## Number of Fisher Scoring iterations: 5

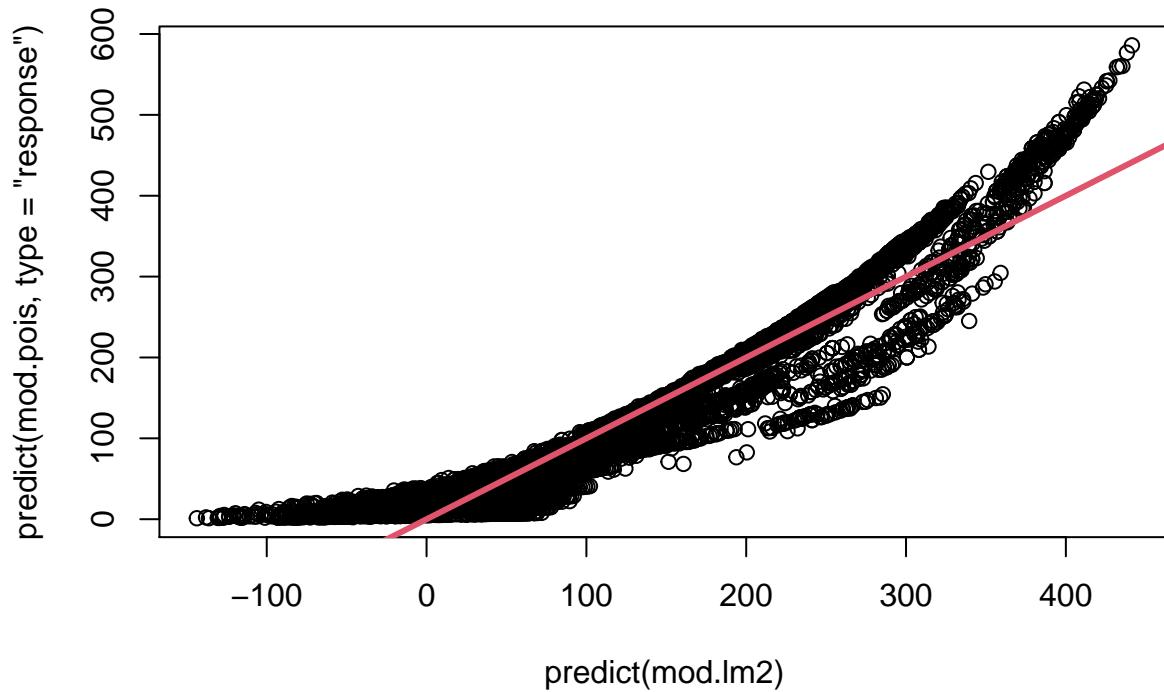
coef.months = c(coef(mod.pois)[2:12], -sum(coef(mod.pois)[2:12]))
plot(coef.months, xlab="Month", ylab="Coeff", xaxt="n", col="blue", pch=19, type="o")
axis(side=1, at=1:12, labels=c("J","F","M","A","M","J","J","A","S","O","N","D"))

coef.hrs = c(coef(mod.pois)[13:35], -sum(coef(mod.pois)[13:35]))
plot(coef.hrs, xlab="Hour", ylab="Coeff", xaxt="n", col="blue", pch=19, type="o")

```



```
# Predict
par(mfrow=c(1,1))
plot(predict(mod.lm2), predict(mod.pois, type="response"))
abline(0,1, col=2, lwd=3)
```



Add a new chunk by clicking the *Insert Chunk* button on the toolbar or by pressing *Ctrl+Alt+I*.