

7 Non-Linear Models

By: Udit (based on ISLR)

Setup

```
library(ISLR2)
library(splines)  # using splines
library(gam)      # using GAM models
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.20
```

```
library(akima)    # surface plots
```

```
attach(Wage)
dim(Wage)
```

```
## [1] 3000  11
```

Polynomial Regression

Regression of Wage ~ Age, up to polynomial of power 4. Polynomial vs. Spline - most important drawback of polynomial being non-locality. That is the fitted function at a given value x_0 depends on data values far from that point.

```
fit.poly <- lm(wage~poly(age,4), data=Wage)
summary(fit.poly)
```

```
##
## Call:
## lm(formula = wage ~ poly(age, 4), data = Wage)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -98.707 -24.626  -4.993  15.217  203.693
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   111.7036     0.7287  153.283 < 2e-16 ***
## poly(age, 4)1   447.0679    39.9148   11.201 < 2e-16 ***
## poly(age, 4)2 -478.3158    39.9148  -11.983 < 2e-16 ***
## poly(age, 4)3   125.5217    39.9148    3.145  0.00168 **
## poly(age, 4)4  -77.9112    39.9148   -1.952  0.05104 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.91 on 2995 degrees of freedom
## Multiple R-squared:  0.08626,    Adjusted R-squared:  0.08504
## F-statistic: 70.69 on 4 and 2995 DF,  p-value: < 2.2e-16
```

```
#Raw polynomials .. poly(age,4, raw=TRUE) or:
fit.temp = lm(wage~age+I(age^2)+I(age^3)+I(age^4),data=Wage) #I - treat it as-is
summary(fit.temp)
```

```
##
## Call:
## lm(formula = wage ~ age + I(age^2) + I(age^3) + I(age^4), data = Wage)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -98.707 -24.626  -4.993  15.217 203.693
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.842e+02  6.004e+01  -3.067 0.002180 **
## age          2.125e+01  5.887e+00   3.609 0.000312 ***
## I(age^2)     -5.639e-01  2.061e-01  -2.736 0.006261 **
## I(age^3)      6.811e-03  3.066e-03   2.221 0.026398 *
## I(age^4)     -3.204e-05  1.641e-05  -1.952 0.051039 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.91 on 2995 degrees of freedom
## Multiple R-squared:  0.08626,    Adjusted R-squared:  0.08504
## F-statistic: 70.69 on 4 and 2995 DF,  p-value: < 2.2e-16
```

Function **poly()** generates a basis of *orthogonal polynomials*, which is preferred. With orthogonal polynomials we can separately test each coefficient. In this case power-4 coefficient is not significant.


References:

- * Visualizing orthogonal polynomials link
- * Raw vs. Orthogonal link

```
#Poly() produces an orthogonal set of basis functions. For example:
par(mfrow=c(1,3))
matplot(seq(-3,3,length.out=90),
        matrix(poly(seq(-3,3,length.out=90),3, raw=FALSE), ncol=3),
        lty = 1, pch = 1, type="l")

matplot(seq(-3,3,length.out=90),
        matrix(poly(seq(-3,3,length.out=90),3, raw=TRUE), ncol=3),
        lty = 1, pch = 1, type="l")

# However the fitted values would be consistent.
plot(fitted(fit.poly), fitted(fit.temp))
```




Plotting the fitted function.

```

agelims=range(age)
age.grid = seq(from =agelims[1], to=agelims[2])
preds = predict(fit.poly, newdata=list(age=age.grid), se=TRUE)
se.bands = cbind(preds$fit+2*preds$se, preds$fit-2*preds$se)

#Plotting
plot(age,wage, col="darkgrey")
lines(age.grid, preds$fit, lwd=2, col="blue") #line-width
matlines(age.grid, se.bands, col="blue", lty=2) #line-type

```



Using `anova()` and **F-test** to test for significance of different variables in a series of nested models.

Fit models ranging from linear to a degree-5 polynomial and seek to determine the simplest model which is sufficient to explain the relationship between wage and age.

Upto Age^3 appears to be significant.

```
fit.1 <- lm(wage~age, data=Wage)
fit.2 <- lm(wage~poly(age,2), data=Wage)
fit.3 <- lm(wage~poly(age,3), data=Wage)
fit.4 <- lm(wage~poly(age,4), data=Wage)
fit.5 <- lm(wage~poly(age,5), data=Wage)
anova(fit.1, fit.2, fit.3, fit.4, fit.5)
```

Analysis of Variance Table

##

Model 1: wage ~ age

Model 2: wage ~ poly(age, 2)

Model 3: wage ~ poly(age, 3)

Model 4: wage ~ poly(age, 4)

Model 5: wage ~ poly(age, 5)

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
--	--------	-----	----	-----------	---	--------

## 1	2998	5022216				
------	------	---------	--	--	--	--

```
## 2    2997 4793430 1    228786 143.5931 < 2.2e-16 ***
## 3    2996 4777674 1    15756   9.8888  0.001679 **
## 4    2995 4771604 1     6070   3.8098  0.051046 .
## 5    2994 4770322 1     1283   0.8050  0.369682
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Since polynomials are orthogonal, we could have simply used the p-values from degree-5 fit to review the results. F-stat is equal to $t - stat^2$

```
coef ( summary (fit.5))
```

```
##              Estimate Std. Error    t value    Pr(>|t|)
## (Intercept)   111.70361   0.7287647  153.2780243 0.000000e+00
## poly(age, 5)1  447.06785  39.9160847   11.2001930 1.491111e-28
## poly(age, 5)2 -478.31581  39.9160847  -11.9830341 2.367734e-32
## poly(age, 5)3  125.52169  39.9160847   3.1446392 1.679213e-03
## poly(age, 5)4  -77.91118  39.9160847  -1.9518743 5.104623e-02
## poly(age, 5)5  -35.81289  39.9160847  -0.8972045 3.696820e-01
```

```
(-11.9830341)^2
```

```
## [1] 143.5931
```

ANOVA method works whether or not we used orthogonal polynomials; it also works when we have other terms in the model as well.

Age^3 is not-significant when other variables are included.

```
fit.a <- lm(wage~education, data=Wage)
fit.b <- lm(wage~education+age, data=Wage)
fit.c <- lm(wage~education+poly(age,2), data=Wage)
fit.d <- lm(wage~education+poly(age,3), data=Wage)
anova(fit.a, fit.b, fit.c, fit.d)
```

```
## Analysis of Variance Table
##
## Model 1: wage ~ education
## Model 2: wage ~ education + age
## Model 3: wage ~ education + poly(age, 2)
## Model 4: wage ~ education + poly(age, 3)
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1     2995 3995721
## 2     2994 3867992  1    127729 102.7378 <2e-16 ***
## 3     2993 3725395  1    142597 114.6969 <2e-16 ***
## 4     2992 3719809  1     5587   4.4936 0.0341 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(fit.a)
```

```
## Analysis of Variance Table
##
## Response: wage
##              Df Sum Sq Mean Sq F value    Pr(>F)
## education      4 1226364   306591   229.81 < 2.2e-16 ***
## Residuals    2995 3995721    1334
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Polynomial - LOGISTIC regression


Change Wage output variable to 0/1, with 1 for >\$250k earners. In GLM due to the way it functions, some of the orthogonality of coefficients is lost, therefore to decide inclusion/exclusion of variable, we'll need to rely on F-test.

Predict() function also provides probabilities, using **type = "response"** option, however that would make the standard-errors/confidence interval non-sensical.

```
fit.log <- glm(I(wage>250)~poly(age,4), data=Wage, family=binomial)
summary(fit.log)
```

```
##
## Call:
## glm(formula = I(wage > 250) ~ poly(age, 4), family = binomial,
##      data = Wage)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.3110  -0.2607  -0.2488  -0.1791   3.7859
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -4.3012     0.3451 -12.465 < 2e-16 ***
## poly(age, 4)1   71.9642    26.1176   2.755  0.00586 **
## poly(age, 4)2  -85.7729    35.9043  -2.389  0.01690 *
## poly(age, 4)3   34.1626    19.6890   1.735  0.08272 .
## poly(age, 4)4  -47.4008    24.0909  -1.968  0.04912 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 730.53  on 2999  degrees of freedom
## Residual deviance: 701.22  on 2995  degrees of freedom
## AIC: 711.22
##
## Number of Fisher Scoring iterations: 9
```


```
plot(age,fitted(fit.log))
```



```
#Calculate Predicted values + Standard Error band for LOGIT!!
preds = predict(fit.log, list(age=age.grid), se=T)
se.bands = preds$fit + cbind(fit=0, lower=-2*preds$se, upper=+2*preds$se)

#Converting from Log-Odds/LOGIT to Probability
prob.bands = exp(se.bands)/(1+exp(se.bands))

#Plotting
matplot(age.grid,prob.bands, col="blue", lwd=c(2,1,1), lty=c(1,2,2), type="l", ylim=c(0,0.1))
points(jitter(age), I(wage>250)/10, pch="|", cex=0.5)
```



```
# fit.log.a <- glm(I(wage>250)~poly(age,2), data=Wage, family=binomial)
# fit.log.b <- glm(I(wage>250)~poly(age,3), data=Wage, family=binomial)
# anova(fit.log.a, fit.log.b)
```

Step Functions

Using `cut()`. Breaks can be manually assigned using **breaks option**.

The age < 33.5 category is left out, so the intercept coefficient of 94 can be interpreted as the average salary for those under 33.5 years of age, and the other coefficients are average additional salary for those other age groups.

```
table(cut(age,4))
```

```
##
## (17.9,33.5] (33.5,49] (49,64.5] (64.5,80.1]
##          750      1399        779         72
```


```
fit.step = lm(wage~cut(age,4), data=Wage)
coef(summary(fit.step))
```

```
##
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) 94.158392 1.476069 63.789970 0.000000e+00
## cut(age, 4)(33.5,49] 24.053491 1.829431 13.148074 1.982315e-38
## cut(age, 4)(49,64.5] 23.664559 2.067958 11.443444 1.040750e-29
## cut(age, 4)(64.5,80.1] 7.640592 4.987424 1.531972 1.256350e-01
```



```
fit.step1 = lm(wage~cut(age,c(17,25,40,65,82)), data=Wage)

plot(age, wage, pch=20, col="gray")
points(age,fitted(fit.step), pch=20, col="blue")
points(age,fitted(fit.step1), col="red")
```



Splines - Fixed-knot Cubic Spline

`bs()` generates the B-spline basis matrix for a polynomial spline (cubic by default.) `ns()` generates natural spline. More explanation [here](#).

```
# 3 knots will lead to 7 DFs (K+4) = 1 intercept + 6 basis functions
# We can either specify knots or DFs
agelims=range(age)
age.grid = seq(from =agelims[1], to=agelims[2])

fit.spline = lm(wage~bs(age, knots=c(25, 40, 60)),data=Wage)
summary(fit.spline)
```

```
##
## Call:
## lm(formula = wage ~ bs(age, knots = c(25, 40, 60)), data = Wage)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -98.832 -24.537  -5.049  15.209  203.207
##
## Coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)                60.494      9.460    6.394 1.86e-10 ***
## bs(age, knots = c(25, 40, 60))1    3.980    12.538    0.317 0.750899
## bs(age, knots = c(25, 40, 60))2    44.631     9.626    4.636 3.70e-06 ***
## bs(age, knots = c(25, 40, 60))3    62.839    10.755    5.843 5.69e-09 ***
## bs(age, knots = c(25, 40, 60))4    55.991    10.706    5.230 1.81e-07 ***
## bs(age, knots = c(25, 40, 60))5    50.688    14.402    3.520 0.000439 ***
## bs(age, knots = c(25, 40, 60))6    16.606    19.126    0.868 0.385338
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.92 on 2993 degrees of freedom
## Multiple R-squared:  0.08642,    Adjusted R-squared:  0.08459
## F-statistic: 47.19 on 6 and 2993 DF,  p-value: < 2.2e-16
```

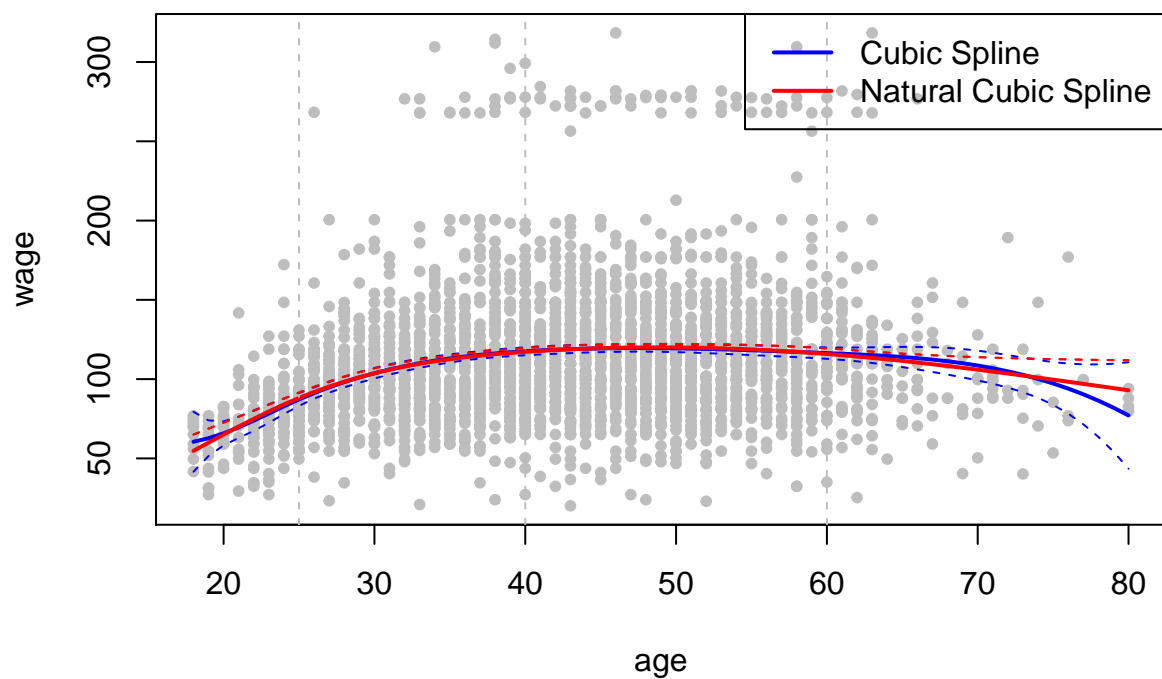
```
pred.spline = predict(fit.spline, list(age=age.grid), se=TRUE)
```

```
fit.nspline = lm(wage~ns(age, knots=c(25, 40, 60)),data=Wage)
summary(fit.spline)
```

```
##
## Call:
## lm(formula = wage ~ bs(age, knots = c(25, 40, 60)), data = Wage)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -98.832 -24.537  -5.049   15.209  203.207
##
## Coefficients:
##                  Estimate Std. Error t value Pr(>|t|)
## (Intercept)          60.494      9.460    6.394 1.86e-10 ***
## bs(age, knots = c(25, 40, 60))1    3.980    12.538    0.317 0.750899
## bs(age, knots = c(25, 40, 60))2    44.631     9.626    4.636 3.70e-06 ***
## bs(age, knots = c(25, 40, 60))3    62.839    10.755    5.843 5.69e-09 ***
## bs(age, knots = c(25, 40, 60))4    55.991    10.706    5.230 1.81e-07 ***
## bs(age, knots = c(25, 40, 60))5    50.688    14.402    3.520 0.000439 ***
## bs(age, knots = c(25, 40, 60))6    16.606    19.126    0.868 0.385338
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.92 on 2993 degrees of freedom
## Multiple R-squared:  0.08642,    Adjusted R-squared:  0.08459
## F-statistic: 47.19 on 6 and 2993 DF,  p-value: < 2.2e-16
```

```
pred.nspline = predict(fit.nspline, list(age=age.grid), se=TRUE)
```

```
plot(age, wage, col="gray", pch=20)
lines(age.grid, pred.spline$fit, col="blue", lwd=2)
lines(age.grid, pred.nspline$fit, col="red", lwd=2)
abline(v=c(25, 40, 60), lty=2, col="grey")
matlines(age.grid, cbind(pred.spline$fit + 2*pred.spline$se,
                          pred.spline$fit - 2*pred.spline$se,
                          pred.nspline$fit + 2*pred.nspline$se,
                          pred.nspline$fit - 2*pred.nspline$se),
          lty=c("solid", "solid", "dashed", "dashed"), col=c("blue", "blue", "red", "red"),
          legend="topright", legend=c("Cubic Spline", "Natural Cubic Spline"), col= c("blue", "red"), lty=1, lwd=2)
```



Splines -

Smoothing-splines


Smoothing-splines do not require knot selection, as **each point is a knot**.
But a smoothing parameter *lambda*. Resulting spline is a natural cubic spline.

```
# Controlling smoothing parameter by (i) Effective degrees of freedom
fit.sm.spline = smooth.spline(age, wage, df=16)
plot(age, wage, col="gray")
lines(fit.sm.spline, col="red", lwd=2)
```

```
# ... or (ii) cross-validation
fit.sm.spline = smooth.spline(age, wage, cv=TRUE)
```

```
## Warning in smooth.spline(age, wage, cv = TRUE): cross-validation with non-unique
## 'x' values seems doubtful
```

```
lines(fit.sm.spline, col="purple", lwd=2)
```



```
fit.sm.spline
```

```
## Call:
## smooth.spline(x = age, y = wage, cv = TRUE)
##
## Smoothing Parameter spar= 0.6988943 lambda= 0.02792303 (12 iterations)
## Equivalent Degrees of Freedom (Df): 6.794596
## Penalized Criterion (RSS): 75215.9
## PRESS(1.o.o. CV): 1593.383
```


LOESS - Local Regression

The larger the span, the smoother the fit.

```
plot(age, wage, xlim=agelims, col="gray")
title("LOESS - Local Regression")

fit1 <- loess(wage~age, span=0.2, data=Wage)
fit2 <- loess(wage~age, span=0.5, data=Wage)
lines(age.grid, predict(fit1, data.frame(age=age.grid)), col="red", lwd=2)
lines(age.grid, predict(fit2, data.frame(age=age.grid)), col="blue", lwd=2)
legend("topright", legend=c("Span=0.2/20%", "Span=0.5/50%"), col=c("Red", "Blue"), lty=1, lwd=2, cex=0.8)
```

LOESS – Local Regression




GAM: Generalized Additive Models

Mixing more than one predictors. Use `s()` to specify a **Smoothing Spline** fit in a GAM Formula.

The generic `plot()` function recognizes that `gam.m3` is an object of class `Gam`, and invokes the appropriate `plot.Gam()` method.

Compelling evidence that a GAM with a linear function of year is better vs. one that doesn't include year (p-value = 0.00014). However, there is no evidence that a non-linear function of year is needed (p-value = 0.349).

```
# Gam with smoothing spline
gam.m3 = gam(wage~s(age,df=5)+s(year,df=4)+education, data=Wage)
par(mfrow=c(1,3))
plot(gam.m3,se=T)
```



Should YEAR be linear or non-linear?

```
gam.m1 <- gam(wage ~ s(age,5) + education, data=Wage)
gam.m2 <- gam(wage ~ year + s(age,5) + education, data=Wage)
anova(gam.m1, gam.m2, gam.m3, test="F")
```

Analysis of Deviance Table

##

Model 1: wage ~ s(age, 5) + education

Model 2: wage ~ year + s(age, 5) + education

Model 3: wage ~ s(age, df = 5) + s(year, df = 4) + education

	Resid. Df	Resid. Dev	Df	Deviance	F	Pr(>F)
## 1	2990	3711731				
## 2	2989	3693842	1	17889.2	14.4771	0.0001447 ***
## 3	2986	3689770	3	4071.1	1.0982	0.3485661

1 2990 3711731

2 2989 3693842 1 17889.2 14.4771 0.0001447 ***

3 2986 3689770 3 4071.1 1.0982 0.3485661

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The “Anova for Parametric Effects” p-values clearly demonstrate that year, age, and education are all highly statistically significant, even when only assuming a linear relationship. Alternatively, the “Anova for Nonparametric Effects” p-values for year and age correspond to a null hypothesis of a linear relationship versus the alternative of a non-linear relationship. The large p-value for year reinforces the conclusion from the ANOVA test that a linear function is adequate for this term.

```
summary(gam.m3)
```

##

Call: gam(formula = wage ~ s(age, df = 5) + s(year, df = 4) + education,
data = Wage)




Deviance Residuals:

##	Min	1Q	Median	3Q	Max
##					

```
## -119.43 -19.70 -3.33 14.17 213.48
##
## (Dispersion Parameter for gaussian family taken to be 1235.69)
##
## Null Deviance: 5222086 on 2999 degrees of freedom
## Residual Deviance: 3689770 on 2986 degrees of freedom
## AIC: 29887.75
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##           Df Sum Sq Mean Sq F value    Pr(>F)
## s(age, df = 5)      1 200684 200684 162.406 < 2.2e-16 ***
## s(year, df = 4)      1  21817  21817  17.655 2.725e-05 ***
## education           4 1069726 267432 216.423 < 2.2e-16 ***
## Residuals          2986 3689770    1236
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##           Npar Df Npar F    Pr(F)
## (Intercept)
## s(age, df = 5)      4 32.380 <2e-16 ***
## s(year, df = 4)      3  1.086 0.3537
## education
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

GAM: LOGIT with smoothing spline



```
gam.l1 = gam(I(wage>250)~s(age,df=4)+year+education, data=Wage, family=binomial)
plot(gam.l1, se=T)
```




#High SE for "<HS Grad" category?





```
##
## education          FALSE TRUE
## 1. < HS Grad        268    0
## 2. HS Grad          966    5
## 3. Some College     643    7
## 4. College Grad     663   22
## 5. Advanced Degree  381   45
```

```
gam.l1.sub = gam(I(wage>250)~s(age,df=4)+year+education, data=Wage,
                 family=binomial,
                 subset=(education != "1. < HS Grad"))
plot(gam.l1.sub, se=T)
```





```
gam.l2.sub = gam(I(wage>250)~s(age,df=4)+s(year,df=4)+education, data=Wage,
  family=binomial,
  subset=(education != "1. < HS Grad"))
plot(gam.l2.sub, se=T)
```





```
anova(gam.l1.sub, gam.l2.sub, test="Chisq") #no-need for adding non-linear terms for 'year'
```


```
## Analysis of Deviance Table
```

```
##
## Model 1: I(wage > 250) ~ s(age, df = 4) + year + education
## Model 2: I(wage > 250) ~ s(age, df = 4) + s(year, df = 4) + education
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1      2723      603.78
## 2      2720      602.87  3   0.90498   0.8242
```

GAM: Using LOESS & Interaction

```
gam.lo <- gam(wage~s(year, df=4) + lo(age, span=0.7) + education, data=Wage)
plot.Gam(gam.lo, se=TRUE, col="green")
```





```
gam.lo.i <- gam(wage~lo(year, age, span=0.5)+ education, data=Wage)
```



```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : liv
## too small. (Discovered by lowesd)
```

```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : lv
## too small. (Discovered by lowesd)
```

```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : liv
## too small. (Discovered by lowesd)
```

```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : lv
## too small. (Discovered by lowesd)
```

```
plot(gam.lo.i)
```



Using GAM plotting functionality with `lm()` models.

```
par(mfrow=c(1,3))
lm1 = lm(wage~ns(age,df=4)+ns(year,df=4)+education,data=Wage)
plot.Gam(lm1, se=T)
```

