

13 Multiple Testing

By: Udit (based on ISLR)

Review of T-test

Simple t-tests using 100 variables, each consisting of 10 observations. The first 50 variables have a non-zero mean of 0.5 by design and variance of 1, while others have mean of 0.

```
set.seed(6)
x = matrix(rnorm(10*100),10,100)
x[,1:50] = x[,1:50] + 0.5
dim(x)
```

```
## [1] 10 100
```

```
t.test(x[,1],mu=0)
```

```
##
## One Sample t-test
##
## data: x[, 1]
## t = 2.0841, df = 9, p-value = 0.06682
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## -0.05171076 1.26242719
## sample estimates:
## mean of x
## 0.6053582
```

```
p.values = rep(0,100)
for(i in 1:100){
  p.values[i] = t.test(x[,i], mu=0)$p.value
}

decision = ifelse(p.values <= 0.05, "Reject Null", "Do Not Reject")
table(decision, true = c(rep("Reject Null",50), rep("Do Not Reject", 50)))
```

```
##           true
## decision  Do Not Reject Reject Null
## Do Not Reject      47      40
## Reject Null       3      10
```

At $\alpha = 0.05$ we reject just 11 out of 50 false null hypotheses. And we would incorrectly reject 3 of the true null hypotheses.

```
# Using Stronger Signal/Noise ratio
x = matrix(rnorm(10*100),10,100)
x[,1:50] = x[,1:50] + 1
p.values = rep(0,100)
for(i in 1:100){
```

```
p.values[i] = t.test(x[,i], mu=0)$p.value
}

decision = ifelse(p.values <= 0.05, "Reject Null", "Do Not Reject")
table(decision, true = c(rep("Reject Null",50), rep("Do Not Reject", 50)))
```


```
##               true
## decision      Do Not Reject Reject Null
## Do Not Reject          49          9
## Reject Null           1         41
```

Family Wise Error Rate (FWER)

```
m = 1:500
fwer1 = 1 - (1-0.05)^m
fwer2 = 1 - (1-0.01)^m
fwer3 = 1 - (1-0.001)^m

matplot(cbind(m,fwer1,fwer2, fwer3), type="l", ylim=c(0,1), col=c(2,3,4), lty=1,
        ylab = "FWER", xlab = "# of Hypotheses",
        main = "FWER: P(Rejecting atleast 1 True Null)")
legend("bottomright", legend=c(0.05, 0.01, 0.001), col=c(2,3,4), lty=1)
```

FWER: P(Rejecting atleast 1 True Null)



FWER - Fund Manager Dataset

```
library(ISLR2)
fund.mini <- Fund[,1:5]
t.test(fund.mini[,1], mu=0)
```

```
##
## One Sample t-test
##
## data: fund.mini[, 1]
## t = 2.8604, df = 49, p-value = 0.006202
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
##  0.8923397 5.1076603
## sample estimates:
## mean of x
##      3
```

```
fund.p = rep(0,5)
for (i in 1:5){
  fund.p[i] = t.test(fund.mini[,i], mu=0)$p.value
}

fund.p
```

```
## [1] 0.006202355 0.918271152 0.011600983 0.600539601 0.755781508
```

```
# Bonferroni adjustment alpha/m
p.adjust(fund.p, method="bonferroni")
```

```
## [1] 0.03101178 1.00000000 0.05800491 1.00000000 1.00000000
```

```
#Holm's adjustment
p.adjust(fund.p, method = "holm")
```

```
## [1] 0.03101178 1.00000000 0.04640393 1.00000000 1.00000000
```

Because the paired t-test below was conducted after visual inspection of the 5 fund managers, in essence, we already carried out 5C2 pairwise comparison through visual inspection. Therefore the p-value should be adjusted for this using **Tukey's HSD (Honest Significant Difference) method**.

```
#Paired t-test
apply(fund.mini, 2, mean)
```

```
## Manager1 Manager2 Manager3 Manager4 Manager5
##      3.0      -0.1       2.8       0.5       0.3
```

```
t.test(fund.mini[,1], fund.mini[,2], paired=T)
```

```
##
## Paired t-test
##
## data: fund.mini[, 1] and fund.mini[, 2]
## t = 2.128, df = 49, p-value = 0.03839
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
```


```
## 0.1725378 6.0274622
## sample estimates:
## mean of the differences
## 3.1

returns = as.vector(as.matrix(fund.mini))
manager = rep(c("1","2","3","4","5"), rep(50,5))
a1 = aov(returns ~ manager) #ANOVA
TukeyHSD(x = a1) #diff between M1 & M2 is no longer significant
```

```
## Tukey multiple comparisons of means
## 95% family-wise confidence level
##
## Fit: aov(formula = returns ~ manager)
##
## $manager
##      diff      lwr      upr    p adj
## 2-1 -3.1 -6.9865435 0.7865435 0.1861585
## 3-1 -0.2 -4.0865435 3.6865435 0.9999095
## 4-1 -2.5 -6.3865435 1.3865435 0.3948292
## 5-1 -2.7 -6.5865435 1.1865435 0.3151702
## 3-2  2.9 -0.9865435 6.7865435 0.2452611
## 4-2  0.6 -3.2865435 4.4865435 0.9932010
## 5-2  0.4 -3.4865435 4.2865435 0.9985924
## 4-3 -2.3 -6.1865435 1.5865435 0.4819994
## 5-3 -2.5 -6.3865435 1.3865435 0.3948292
## 5-4 -0.2 -4.0865435 3.6865435 0.9999095
```

```
plot(TukeyHSD(x = a1))
```

95% family-wise confidence level



Rate (FDR)

False Discovery

Far too many tests to control for FWER (since that would be impossibly punitive and lead to extremely few ‘discoveries’.) Instead, we focus on FDR: expected fraction of rejected null hypotheses that are actually false positives.

```
fund.p = rep(0,2000)
for (i in 1:2000){
  fund.p[i] = t.test(Fund[,i],mu=0)$p.value
}
fund.p[1:5]
```

```
## [1] 0.006202355 0.918271152 0.011600983 0.600539601 0.755781508
```

```
# Benjamini-Hochberg adjustment
q.BH = p.adjust(fund.p, method="BH")
q.BH[1:10] # q.value is the lower FDR at which that H0 can be rejected
```

```
## [1] 0.08988921 0.99149100 0.12211561 0.92342997 0.95603587 0.07513802
## [7] 0.07670150 0.07513802 0.07513802 0.07513802
```

```
# Rejected Nulls for FDR of 10%
sum(q.BH <= 0.1) # 146.. we can expect ~15 of these to be false positives
```


```
## [1] 146
```

```
# Bonferroni would be extremely punitive
sum(fund.p <= (0.1/2000)) # 0 discoveries
```

```
## [1] 0
```

```
# BH method - arrange p-values, compare with q*j/m
m = length(fund.p)
p = sort(fund.p)
q = .1
idx = which(p < q*(1:m)/m)

plot(p, log="xy", ylim=c(4e-6,1), ylab="P-value", xlab="Index", main="", pch=20, col="gray")
points(idx, p[idx], col=4, pch=20)
abline(a=0, b = (q/m), col=2, lty=1)
abline(h=0.1/2000, col=3)
```



Resampling Approach

```
attach(Khan) #Khan gene data
x = rbind(xtrain, xtest)
y = c(ytrain, ytest)
table(y) # Four classes of cancer
```

```
## y
##  1  2  3  4
## 11 29 18 25
```

```
# Comparing 11 gene's difference b/w class 2 and 4
x1 = x[which(y==2),]
x2 = x[which(y==4),]
t.out = t.test(x1[,11], x2[,11], var.equal=TRUE)
t.out$statistic; t.out$p.value # p-value based on 'theoretical' distribution
```

```
##          t
## -2.093633
```

```
## [1] 0.04118644
```

```
# Re-sampling to build empirical distribution
n1 = nrow(x1)
n2 = nrow(x2)

set.seed(1)
```

```

b = 10000
t.b = rep(NA,b)
for(i in 1:b){
  dat = sample(c(x1[,11],x2[,11])) # jumbles up all values
  t.b[i] = t.test(dat[1:n1], dat[(n1+1):(n1+n2)], var=T)$statistic
}
mean(abs(t.b) >= abs(t.out$statistic)) #0.0416.. same as theoretical distribution

```


```
## [1] 0.0416
```

```

hist(t.b, breaks=100, xlim=c(-4.2, 4.2), xlab="Null Distribution of Test Stat")
lines(seq(-4.2,4.2,len=1000),dt(seq(-4.2,4.2,len=1000), df=(n1+n2-2))*1000)
text(t.out$statistic, 350, paste("T = ", round(t.out$statistic,4), sep=""))

```

Histogram of t.b



Calculating FDR

for all 2,308 genes.

```

m = 50 #taking 50 genes at random
set.seed(1)
index = sample(ncol(x), m)
Ts = rep(NA, m)
Ts.star = matrix(NA, ncol = m, nrow = b)
for(j in 1:m){
  k = index[j]
  Ts[j] = t.test(x1[,k], x2[,k], var.equal=TRUE)$statistic
  for (i in 1:b){
    dat = sample(c(x1[,k],x1[,k]))
    Ts.star[i,j] = t.test(dat[1:n1], dat[(n1+1):(n1+n2)], var.equal=TRUE)$statistic
  }
}

```

```

cs = sort(abs(Ts))
FDRs = Rs = Vs = rep(NA,m)
for(j in 1:m){
  Vs[j] = sum(abs(Ts.star) >= cs[j])/b      # V = false rejections
  Rs[j] = sum(abs(Ts) >= cs[j])              # R = total rejected nulls.. 50:1
  FDRs[j] = Vs[j]/Rs[j]
}

max(Rs[FDRs <= 0.1]) #6 out of 50 nulls can be rejected.. expect ~1 false positive

```

```
## [1] 6
```

```
max(Rs[FDRs <= 0.3]) #15 out of 50 nulls can be rejected.. expect ~2 false positive
```

```
## [1] 15
```

```
plot(Rs, FDRs, xlab="Number of Rejections", type="l", ylab="FDR", col=4, lwd=3)
```

